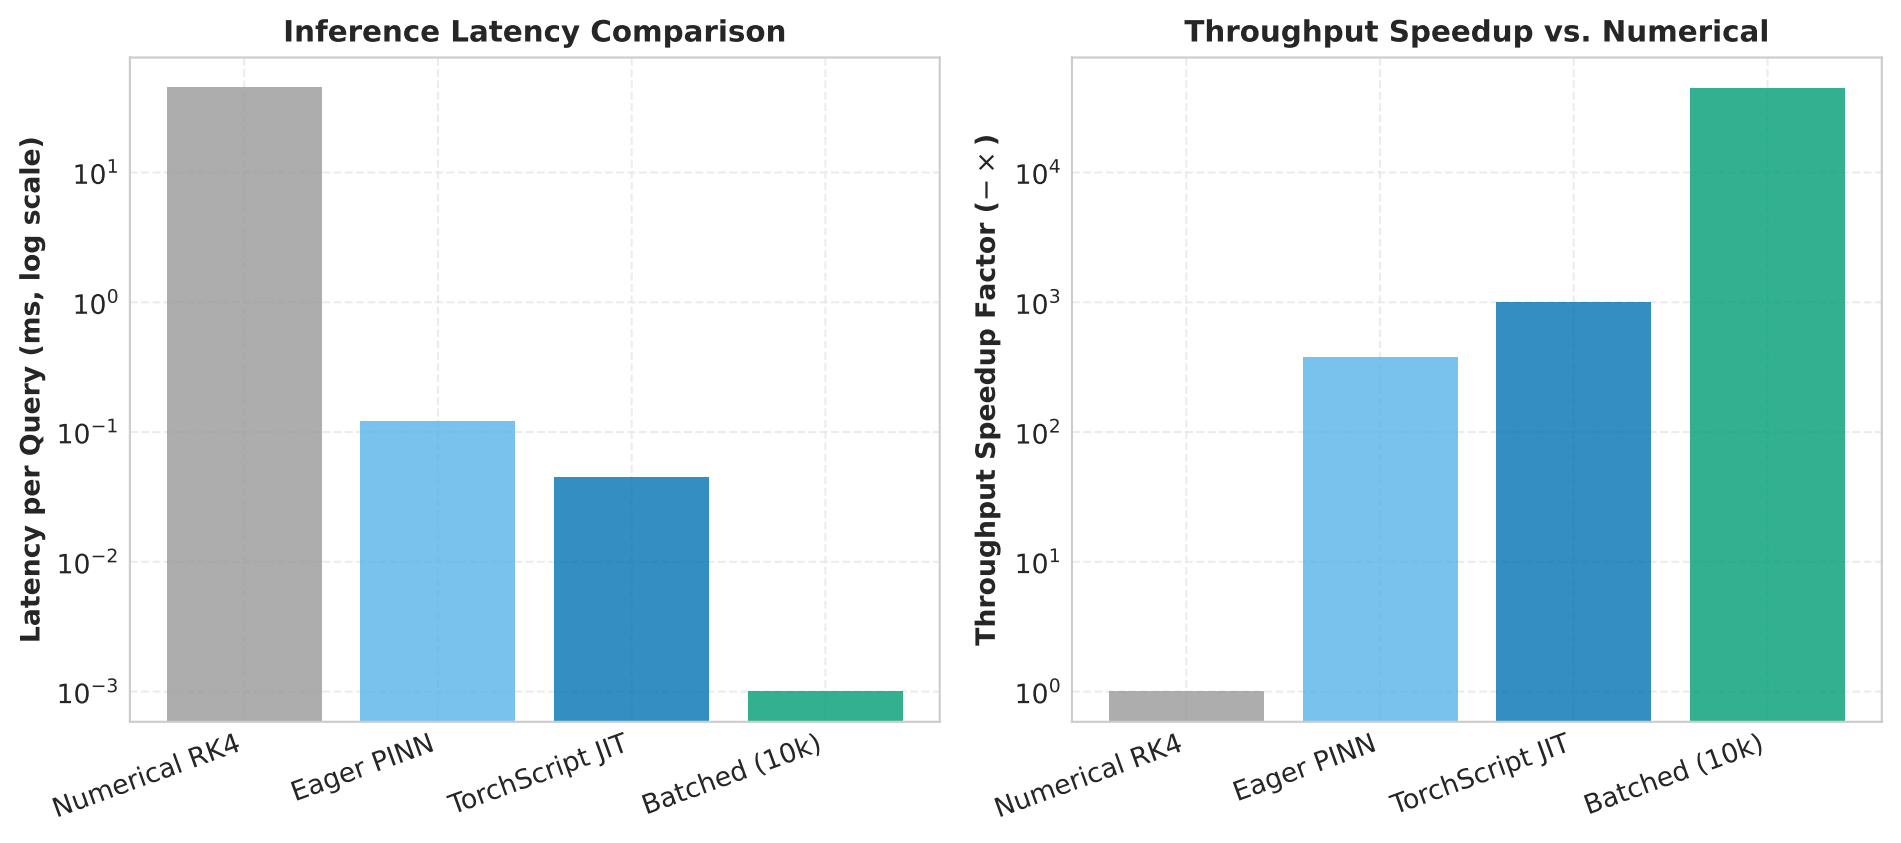
Inference Latency Comparison
Latency per Query (ms, log scale)
10^{1}
10^{0}
10^{-1}
10^{-2}
10^{-3}
Numerical RK4
Eager PINN
TorchScript JIT
Batched (10k)
Throughput Speedup vs. Numerical
Throughput Speedup Factor (− ×)
10^{4}
10^{3}
10^{2}
10^{1}
10^{0}
Numerical RK4
Eager PINN
TorchScript JIT
Batched (10k)